CREATE TABLES IN HIVE AND EXECUTE QUERIES

hadoop@vinoth-ubuntu:~$ ls /usr/share/java
java-atk-wrapper.jar java_defaults.mk libintl-0.21.jar libintl.jar
hadoop@vinoth-ubuntu:~$ wget
https://dev.mysql.com/get/Downloads/Connector-J/mysql-connector-j-9.0.0.tar.gz
--2024-09-13 12:08:52--
https://dev.mysql.com/get/Downloads/Connector-J/mysql-connector-j-9.0.0.tar.gz
Resolving dev.mysql.com (dev.mysql.com)... 23.58.39.183, 2600:140f:5e00:199::2e31, 2600:140f:5e00:18b::2e31
Connecting to dev.mysql.com (dev.mysql.com)|23.58.39.183|:443... connected.
HTTP request sent, awaiting response... 302 Moved Temporarily
Location: https://cdn.mysql.com//Downloads/Connector-J/mysql-connector-j-9.0.0.tar.gz [following]
--2024-09-13 12:08:53--
https://cdn.mysql.com//Downloads/Connector-J/mysql-connector-j-9.0.0.tar.gz
Resolving cdn.mysql.com (cdn.mysql.com)... 23.58.42.45, 2600:140f:5e00:187::1d68, 2600:140f:5e00:19e::1d68
Connecting to cdn.mysql.com (cdn.mysql.com)|23.58.42.45|:443... connected.
HTTP request sent, awaiting response... 200 OK
Length: 4469329 (4.3M) [application/x-tar-gz]
Saving to: ‘mysql-connector-j-9.0.0.tar.gz’

mysql-connector-j-9.0 100%[=====================>] 4.26M 6.57MB/s in 0.6s

2024-09-13 12:08:54 (6.57 MB/s) - ‘mysql-connector-j-9.0.0.tar.gz’ saved [4469329/4469329]

hadoop@vinoth-ubuntu:~$ tar -xvzf mysql-connector-j-9.0.0.tar.gz
hadoop@vinoth-ubuntu:~$ ls
apache-hive-3.1.3-bin hadoop-3.3.6.tar.gz pig
apache-hive-3.1.3-bin.tar.gz hadoopdata pig-0.16.0.tar.gz
derby.log Music Public
Desktop mysql-connector-j-9.0.0 snap
Documents mysql-connector-j-9.0.0.tar.gz Templates
Downloads mysql-connector-java-8.0.15.jar Videos
hadoop Pictures
hadoop@vinoth-ubuntu:~$ cd mysql-connector-j-9.0.0
hadoop@vinoth-ubuntu:~/mysql-connector-j-9.0.0$ ls
build.xml INFO_BIN LICENSE README

CHANGES INFO_SRC mysql-connector-j-9.0.0.jar src
hadoop@vinoth-ubuntu:~/mysql-connector-j-9.0.0$ sudo cp mysql-connector-j-9.0.0.jar /usr/share/java/mysql-connector.java.jar
[sudo] password for hadoop:
hadoop@vinoth-ubuntu:~/mysql-connector-j-9.0.0$ sudo cp mysql-connector-j-9.0.0.jar /home/hadoop/apache-hive-3.1.3-bin/lib/mysql-connector-java.jar
hadoop@vinoth-ubuntu:~/mysql-connector-j-9.0.0$ ls $HIVE_HOME/bin/schematool
/home/hadoop/apache-hive-3.1.3-bin/bin/schematool
hadoop@vinoth-ubuntu:~/mysql-connector-j-9.0.0$
hadoop@vinoth-ubuntu:~/mysql-connector-j-9.0.0$ schematool -initSchema -dbType mysql --verbose
No rows affected (0.004 seconds)
0: jdbc:mysql://localhost/metastore> !closeall
Closing: 0: jdbc:mysql://localhost/metastore?createDatabaseIfNotExist=true
beeline>
beeline> Initialization script completed
schemaTool completed
hadoop@vinoth-ubuntu:~/mysql-connector-j-9.0.0$ hive
SLF4J: Class path contains multiple SLF4J bindings.
SLF4J: Found binding in [jar:file:/home/hadoop/apache-hive-3.1.3-bin/lib/log4j-slf4j-impl-2.17.1.jar!/org/slf4j/impl/StaticLoggerBinder.class]
SLF4J: Found binding in [jar:file:/home/hadoop/hadoop/share/hadoop/common/lib/slf4j-reload4j-1.7.36.jar!/org/slf4j/impl/StaticLoggerBinder.class]
SLF4J: See http://www.slf4j.org/codes.html#multiple_bindings for an explanation.
SLF4J: Actual binding is of type [org.apache.logging.slf4j.Log4jLoggerFactory]
Hive Session ID = 678d5628-cbdd-4c89-93ec-b44e8ccd61ce

Logging initialized using configuration in jar:file:/home/hadoop/apache-hive-3.1.3-bin/lib/hive-common-3.1.3.jar!/hive-log4j2.properties Async: true
Hive-on-MR is deprecated in Hive 2 and may not be available in the future versions. Consider using a different execution engine (i.e. spark, tez) or using Hive **1.X releases.**
Hive Session ID = 16b26a3d-b4b5-4a47-80c8-803f7e82a0ac
hive> show databases;
OK
default
Time taken: 0.559 seconds, Fetched: 1 row(s)
hive>

5A HIVE COMMANDS

hadoop@vinoth-ubuntu:~$ hive
SLF4J: Class path contains multiple SLF4J bindings.
SLF4J: Found binding in
[jar:file:/home/hadoop/apache-hive-3.1.3-bin/lib/log4j-slf4j-impl-2.17.1.jar!/org/slf4j/impl/StaticLoggerBinder.class]
SLF4J: Found binding in
[jar:file:/home/hadoop/hadoop/share/hadoop/common/lib/slf4j-reload4j-1.7.36.jar!/org/slf4j/impl/StaticLoggerBinder.class]
SLF4J: See http://www.slf4j.org/codes.html#multiple_bindings for an explanation.
SLF4J: Actual binding is of type [org.apache.logging.slf4j.Log4jLoggerFactory]
Hive Session ID = 75d9d072-33d3-4769-b257-150080dab6bb

Logging initialized using configuration in
jar:file:/home/hadoop/apache-hive-3.1.3-bin/lib/hive-common-3.1.3.jar!/hive-log4j2.properties Async: true
Hive-on-MR is deprecated in Hive 2 and may not be available in the future versions. Consider using a different execution engine (i.e. spark, tez) or using Hive 1.X releases.
Hive Session ID = d2b98337-9fb3-4319-bdb3-321c9305ac3f
hive> create database financials;
OK
Time taken: 0.472 seconds
hive> use financials;
OK
Time taken: 0.085 seconds
hive> create table finance_table(id int, name String);
OK
Time taken: 0.962 seconds
hive> **insert into finance_table values (1,'Alice'),(2,'Bob'),(3,'Charlie');**
Query ID = hadoop_20240913125807_4b8ec766-dc28-44f6-ab0a-1253fa7d3e1c
Total jobs = 3
Launching Job 1 out of 3
Number of reduce tasks determined at compile time: 1
In order to change the average load for a reducer (in bytes):
  set hive.exec.reducers.bytes.per.reducer=<number>
In order to limit the maximum number of reducers:
  set hive.exec.reducers.max=<number>
In order to set a constant number of reducers:
  set mapreduce.job.reduces=<number>
Job running in-process (local Hadoop)
2024-09-13 12:58:12,694 Stage-1 map = 0%, reduce = 0%

2024-09-13 12:58:16,005 Stage-1 map = 100%, reduce = 100%
Ended Job = job_local504772510_0001
Stage-4 is selected by condition resolver.
Stage-3 is filtered out by condition resolver.
Stage-5 is filtered out by condition resolver.
Moving data to directory
hdfs://localhost:9000/user/hive/warehouse/financials.db/finance_table/.hive-staging_hive_2024-09-13_12-58-07_881_6581065941079588430-1/-ext-10000
Loading data to table financials.finance_table
MapReduce Jobs Launched:
Stage-Stage-1: HDFS Read: 0 HDFS Write: 208 SUCCESS
Total MapReduce CPU Time Spent: 0 msec
OK
Time taken: 10.489 seconds
hive> create view myview as select name.id from finace_table;
FAILED: SemanticException [Error 10001]: Line 1:42 Table not found 'finace_table'
hive> create view myview as select name.id from finance_table;
FAILED: SemanticException [Error 10042]: Line 1:29 . Operator is only supported on struct or list of struct types 'id'
hive> create view myview as select name,id from finance_table;
OK
Time taken: 0.328 seconds
hive> select * from myview;
OK
Alice 1
Bob   2
Charlie   3
Time taken: 0.249 seconds, Fetched: 3 row(s)
hive> describe finance_table;
OK
id   int
name   string
Time taken: 0.108 seconds, Fetched: 2 row(s)
hive> alter table finance_table add columns(age int);
OK
Time taken: 0.242 seconds
hive> describe finance_table;
OK
id   int
name   string
age   int
Time taken: 0.088 seconds, Fetched: 3 row(s)
hive> quit;