

For hadoop.3x version

<http://localhost:9870>

<http://localhost:8088/cluster>

Place the mapper file ,reducer file and hadoop dtarming jar file in Documents,Create input foler in hadoop and place the wordcount.txt file on it.

hadoop/Documents\$ give below comments to run

```
hadoop@Ubuntu:~/Documents$ hadoop jar hadoop-streaming-2.7.3.jar -input
/home/hadoop/input/word_count_data.txt -output /home/hadoop/output -mapper mapper.py -
reducer reducer.py
```

To check the output folder part-oooo file is created or not

```
hadoop@Ubuntu:~/Documents$ hadoop fs -ls /home/hadoop/output
```

```
hadoop@Ubuntu:~/Documents$ hadoop fs -ls /home/hadoop/output
```

Found 2 items

```
-rw-r--r--  1 hadoop supergroup      0 2024-08-03 08:59 /home/hadoop/output/_SUCCESS
-rw-r--r--  1 hadoop supergroup    592 2024-08-03 08:59 /home/hadoop/output/part-00000
```

```
hadoop@Ubuntu:~/Documents$ hdfs dfs -cat /user/hadoop/output/part-00000
```

cat: `/user/hadoop/output/part-00000': No such file or directory

### Verify the output

```
hadoop@Ubuntu:~/Documents$ hdfs dfs -cat /home/hadoop/output/part-00000
```

```
2,000 1
ChatGPT 1
Did 1
Roman2
Romans 1
Some 1
Sure! 1
This 1
a 3
actually 1
ancient 1
and 3
ash 1
ash, 1
because 1
buildings 1
called 1
concrete 2
concrete, 1
```

concrete.	1
construction	1
durable	1
for	1
form	2
from	1
gets	1
harbors	1
has	1
have	1
impressive	1
incredibly	1
is	1
know	1
lime,	1
longevity	1
made	1
many	1
mineral	1
mixture	1
modern	1
of	3
outlasted	1
over	2
partly	1
reacts	1
reinforces	1
seawater	1
seawater,	1
showcasing	1
stronger	1
structures?	1
survived	1
techniques	1
that	2
the	5
their	1
time.	1
to	1
tobermorite,	1
used	1
volcanic	2
was	1
which	1
with	1
years,	1
you	1

path of hadoop input file  
 /home/hadoop/input/word\_count\_data.txt

path of hadoop output file

/home/hadoop/output

---

Commands on hadoop to check the input and output file

## 1. List Contents of a Hadoop Directory

To list the contents of a directory in HDFS, use the `-ls` option with the `hadoop fs` command.

List the Input Folder

```
bash
```

```
hadoop fs -ls /user/hduser/input
```

List the Output Folder

```
bash
```

```
hadoop fs -ls /user/hduser/output
```

## 2. View Detailed Information

The `-ls` command provides a detailed listing of files and directories, including permissions, owner, group, size, and modification date.

## 3. View File Contents

To view the contents of a file, use the `-cat` option. For example:

View a File in the Input Folder

```
bash
```

```
hadoop fs -cat /user/hduser/input/filename.txt
```

Replace filename.txt with the actual name of the file you want to view.

View a File in the Output Folder

If your output folder contains multiple files (e.g., part-r-000000), you can view one of the files:

```
bash
```

```
hadoop fs -cat /user/hduser/output/part-r-000000
```

#### 4. Check for Folder Existence

To check if a folder exists in HDFS, you can use the `-test` command with the `-d` option:

Check if Input Folder Exists

```
bash
```

```
hadoop fs -test -d /user/hduser/input && echo "Input folder exists" || echo "Input folder does not exist"
```

Check if Output Folder Exists

```
bash
```

```
hadoop fs -test -d /user/hduser/output && echo "Output folder exists" || echo "Output folder does not exist"
```

#### 5. Additional Commands

Make a Directory:

```
bash
```

```
hadoop fs -mkdir /user/hduser/new_directory
```

Delete a File or Directory:

```
bash
```

```
hadoop fs -rm /user/hduser/input/filename.txt
```

```
hadoop fs -rm -r /user/hduser/output
```

Get File Status:

```
bash
```

```
hadoop fs -stat /user/hduser/input/filename.txt
```

Example Usage

Assuming you have the following setup:

Input Folder: /user/hduser/input

Output Folder: /user/hduser/output

You would list the contents like this:

```
bash
```

```
hadoop fs -ls /user/hduser/input
```

And to check the output:

```
bash
```

```
hadoop fs -ls /user/hduser/output
```

```
bash
```

```
hadoop fs -cat /user/hduser/output/part-r-00000
```

---

## PIG UDF PROGRAM

To check the pig program

---

```
hadoop@Ubuntu:~/Documents$ nano sample.txt
```

Paste the below content to sample.txt

```
1,John
```

```
2,Jane
```

```
3,Joe
```

```
4,Emma
```

```
hadoop@Ubuntu:~/Documents$ hadoop fs -put sample.txt /home/hadoop/piginput/
```

---

```
hadoop@Ubuntu:~/Documents$ nano demo_pig.pig
```

paste the below the content to demo\_pig.pig

```
-- Load the data from HDFS
```

```
data = LOAD '/home/hadoop/piginput/sample.txt' USING PigStorage(',') AS (id:int>
```

```
-- Dump the data to check if it was loaded correctly
```

```
DUMP data;
```

---

```
hadoop@Ubuntu:~/Documents$ pig demo_pig.pig
```

2024-08-07 12:13:08,791 [main] INFO

org.apache.pig.backend.hadoop.executionengine.util.MapRedUtil - Total input paths to process : 1

(1,John)

(2,Jane)

(3,Joe)

(4,Emma)

By using these commands, you can manage and inspect files and directories in your Hadoop setup.

-----up-----

To Run pig basic program and uf program

uppercase\_udf.py

-----

```
def uppercase(text):
```

```
    return text.upper()
```

```
if __name__ == "__main__":
```

```
    import sys
```

```
    for line in sys.stdin:
```

```
        line = line.strip()
```

```
        result = uppercase(line)
```

```
        print(result)
```

-----

Create the udfs folder on hadoop

```
hadoop@Ubuntu:~/Documents$ hadoop fs -mkdir /home/hadoop/udfs
```

put the upppercase\_udf.py in to the abv folder

```
hadoop@Ubuntu:~/Documents$ hdfs dfs -put uppercase_udf.py /home/hadoop/udfs/
```

```
hadoop@Ubuntu:~/Documents$ nano udf_example.pig
```

### **udf\_example.pig**

```
-----  
-- Register the Python UDF script  
REGISTER 'hdfs:///home/hadoop/udfs/uppercase_udf.py' USING jython AS udf;  
  
-- Load some data  
data = LOAD 'hdfs:///home/hadoop/sample.txt' AS (text:chararray);  
  
-- Use the Python UDF  
uppercased_data = FOREACH data GENERATE udf.uppercase(text) AS uppercase_text;  
  
-- Store the result  
STORE uppercased_data INTO 'hdfs:///home/hadoop/pig_output_data';  
-----
```

place sample.txt file on hadoop

```
hadoop@Ubuntu:~/Documents$ hadoop fs -put sample.txt /home/hadoop/
```

To Run the pig file

```
hadoop@Ubuntu:~/Documents$ pig -f udf_example.pig
```

finally u get

Success!

Job Stats (time in seconds):

JobId	Maps	Reduces	MaxMapTime	MinMapTime	AvgMapTime	MedianMapTime	MaxReduceTime	MinReduceTime	AvgReduceTime	MedianReductime
	Alias	Feature	Outputs							
job_local1786848041_0001	1	0	n/a	n/a	n/a	n/a	00	0	0	
	data,uppercased_data	MAP_ONLY	hdfs:///home/hadoop/pig_output_data,							



Input(s):

Successfully read 4 records (42778068 bytes) from: "hdfs:///home/hadoop/sample.txt"

Output(s):

Successfully stored 4 records (42777870 bytes) in: "hdfs:///home/hadoop/pig\_output\_data"

Counters:

Total records written : 4

Total bytes written : 42777870

Spillable Memory Manager spill count : 0

Total bags proactively spilled: 0

Total records proactively spilled: 0

Job DAG:

job\_local1786848041\_0001

2024-08-07 13:33:04,631 [main] WARN org.apache.hadoop.metrics2.impl.MetricsSystemImpl - JobTracker metrics system already initialized!

2024-08-07 13:33:04,639 [main] WARN org.apache.hadoop.metrics2.impl.MetricsSystemImpl - JobTracker metrics system already initialized!

2024-08-07 13:33:04,644 [main] WARN org.apache.hadoop.metrics2.impl.MetricsSystemImpl - JobTracker metrics system already initialized!

2024-08-07 13:33:04,667 [main] INFO

org.apache.pig.backend.hadoop.executionengine.mapReduceLayer.MapReduceLauncher - Success!

To check the output file is created

hadoop@Ubuntu:~/Documents\$ hdfs dfs -ls /home/hadoop/pig\_output\_data

Found 2 items

If you need to examine the files in the output folder, use:

To view the output

hadoop@Ubuntu:~/Documents\$ hdfs dfs -cat /home/hadoop/pig\_output\_data/part-m-000000

1,JOHN

2,JANE

3,JOE

4,EMMA

-----

Create json file on bash & save as emp.json

nano emp.json ; Paste the below content on it

```
[
  {"name": "John Doe", "age": 30, "department": "HR", "salary": 50000},
  {"name": "Jane Smith", "age": 25, "department": "IT", "salary": 60000},
  {"name": "Alice Johnson", "age": 35, "department": "Finance", "salary": 70000},
  {"name": "Bob Brown", "age": 28, "department": "Marketing", "salary": 55000},
  {"name": "Charlie Black", "age": 45, "department": "IT", "salary": 80000}
]
```

Check json is readable or any error by giving

**install jq by sudo apt-get install jq**

**hadoop@Ubuntu:~\$ jq . emp.json**

```
[
  {
    "name": "John Doe",
    "age": 30,
    "department": "HR",
    "salary": 50000
  },
  {
    "name": "Jane Smith",
    "age": 25,
    "department": "IT",
    "salary": 60000
  },
  {
    "name": "Alice Johnson",
```

```

    "age": 35,
    "department": "Finance",
    "salary": 70000
  },
  {
    "name": "Bob Brown",
    "age": 28,
    "department": "Marketing",
    "salary": 55000
  },
  {
    "name": "Charlie Black",
    "age": 45,
    "department": "IT",
    "salary": 80000
  }
]

```

**bash:** put the employees.json local directory to *home/hadoop* directory

## Example

Suppose the original employees relation has the following data:

name	age	department	salary
John Doe	30	HR	50000
Jane Smith	25	IT	60000
Alice Johnson	35	Finance	70000
Bob Brown	28	Marketing	55000
Charlie Black	45	IT	80000

After executing:

pig shell: Load the json file by giving following command

```

grunt>-- Load the data employees = LOAD '/home/hadoop/emp.json' USING
JsonLoader('name:chararray,age:int,department:chararray,salary:float');

```

```

grunt>projected = FOREACH employees GENERATE name, salary;

```

DUMP projected;

The projected relation will look like:

<b>name</b>	<b>salary</b>
John Doe	50000
Jane Smith	60000
Alice Johnson	70000
Bob Brown	55000
Charlie Black	80000

-----

Assume your employees dataset looks like this:

name	age	department	salary
John Doe	30	HR	50000
Jane Smith	25	IT	60000
Alice Johnson	35	Finance	70000
Bob Brown	28	Marketing	55000
Charlie Black	45	IT	80000

#### 1. Aggregation

Aggregate the total salary:

pig

-- Load the data

```
employees = LOAD '/home/hadoop/employees.json' USING  
JsonLoader('name:chararray,age:int,department:chararray,salary:float');
```

-- Aggregate: Calculate the total salary

```
total_salary = FOREACH (GROUP employees ALL) GENERATE SUM(employees.salary) AS  
total_salary;
```

DUMP total\_salary;

Output:

scss

(315000.0)

## 2. Skip

Skip the first 2 records:

pig

-- Load the data

```
employees = LOAD '/home/hadoop/employees.json' USING  
JsonLoader('name:chararray,age:int,department:chararray,salary:float');
```

-- Skip the first 2 records

```
skipped_employees = LIMIT employees 1000000; -- Use LIMIT to handle skipping
```

```
DUMP skipped_employees;
```

Output:

name	age	department	salary
Alice Johnson	35	Finance	70000
Bob Brown	28	Marketing	55000
Charlie Black	45	IT	80000

Note: The LIMIT command should be used with an appropriate number, as Pig does not directly support skipping a specific number of records.

## 3. Limit

Limit the results to the top 3 records:

pig

-- Load the data

```
employees = LOAD '/home/hadoop/employees.json' USING  
JsonLoader('name:chararray,age:int,department:chararray,salary:float');
```

-- Limit: Get the top 3 highest earners

```
top_3_employees = LIMIT employees 3;
```

```
DUMP top_3_employees;
```

Output:

name	age	department	salary
Charlie Black	45	IT	80000
Alice Johnson	35	Finance	70000
Jane Smith	25	IT	60000

#### 4. Count

Count the number of employees:

pig

-- Load the data

```
employees = LOAD '/home/hadoop/employees.json' USING  
JsonLoader('name:chararray,age:int,department:chararray,salary:float');
```

-- Count the number of employees

```
employee_count = FOREACH (GROUP employees ALL) GENERATE COUNT(employees) AS  
total_count;
```

```
DUMP employee_count;
```

Output:

scss

(5)

## 5. Remove

Remove employees from a specific department, e.g., "IT":

pig

-- Load the data

```
employees = LOAD '/home/hadoop/employees.json' USING  
JsonLoader('name:chararray,age:int,department:chararray,salary:float');
```

-- Remove employees from the 'IT' department

```
filtered_employees = FILTER employees BY department != 'IT';
```

```
DUMP filtered_employees;
```

Output:

name	age	department	salary
John Doe	30	HR	50000
Alice Johnson	35	Finance	70000
Bob Brown	28	Marketing	55000

=====

**import Json file and do projection, aggregation, limit,count ,skip and remove using python and hdfs**

Steps to be followed:

**Install** pandas and hdfs using pip.

- **Optionally** install pyarrow or hdfs3 if needed based on your specific requirements.
- **Verify** the installation to ensure everything is set up correctly.

## Required Packages

pandas:

Purpose: Provides data structures and functions to efficiently manipulate and analyze data.

Installation: Use pip to install pandas.

```
bash
```

```
pip install pandas
```

hdfs:

Purpose: Provides a Python interface to interact with HDFS.

Installation: Use pip to install hdfs.

```
bash
```

```
pip install hdfs
```

## Additional Considerations

While the script should work with just the above packages, here are some additional considerations:

pyarrow (Optional but useful):

Purpose: If you're working with Apache Arrow or need additional features for handling large datasets or different file formats, pyarrow can be useful.

Installation: Use pip to install pyarrow.



```
bash
```

```
pip install pyarrow
```

hdfs3 (Alternative to hdfs):

Purpose: Another Python library for interacting with HDFS. It's an alternative to the hdfs package and might be preferred in some scenarios.

Installation: Use pip to install hdfs3.

```
bash
```

```
pip install hdfs3
```

### Verifying Package Installation

After installing the required packages, you can verify that they are correctly installed and accessible in your Python environment:

```
python
```

```
import pandas as pd
```

```
from hdfs import InsecureClient
```

```
# Check pandas version
```

```
print("Pandas version:", pd.__version__)
```

```
# Test HDFS client connection
```

```
client = InsecureClient('http://localhost:9870', user='hadoop')
```

```
print("HDFS status:", client.status('/'))
```

If you run this script and see the version of pandas and a status message from HDFS without any errors, the packages are installed correctly.

Create process\_data.py file

```
from hdfs import InsecureClient
```

```
import pandas as pd
```

```
import json
```

```
# Connect to HDFS
```

```
hdfs_client = InsecureClient('http://localhost:9870', user='hdfs')
```

```
# Read JSON data from HDFS
```

```
try:
```

```
    with hdfs_client.read('/home/hadoop/emp.json', encoding='utf-8') as reader:
```

```
        json_data = reader.read() # Read the raw data as a string
```

```
        if not json_data.strip(): # Check if data is empty
```

```
            raise ValueError("The JSON file is empty.")
```

```
        print(f"Raw JSON Data: {json_data[:1000]}") # Print first 1000 characters for debugging
```

```
        data = json.loads(json_data) # Load the JSON data
```

```
except json.JSONDecodeError as e:
```

```
    print(f"JSON Decode Error: {e}")
```

```
    exit(1)
```

```
except Exception as e:
```

```
    print(f"Error reading or parsing JSON data: {e}")
```

```
    exit(1)
```

```
# Convert JSON data to DataFrame
```

```
try:
```

```
    df = pd.DataFrame(data)
```

```
except ValueError as e:
```

```
    print(f"Error converting JSON data to DataFrame: {e}")
```

```
    exit(1)
```

```
# Projection: Select only 'name' and 'salary' columns
```

```
projected_df = df[['name', 'salary']]

# Aggregation: Calculate total salary
total_salary = df['salary'].sum()

# Count: Number of employees earning more than 50000
high_earners_count = df[df['salary'] > 50000].shape[0]

# Limit: Get the top 5 highest earners
top_5_earners = df.nlargest(5, 'salary')

# Skip: Skip the first 2 employees
skipped_df = df.iloc[2:]

# Remove: Remove employees from a specific department
filtered_df = df[df['department'] != 'IT']

# Save the filtered result back to HDFS
filtered_json = filtered_df.to_json(orient='records')

try:
    with hdfs_client.write('/home/hadoop/filtered_employees.json', encoding='utf-8',
        overwrite=True) as writer:
        writer.write(filtered_json)
    print("Filtered JSON file saved successfully.")
except Exception as e:
    print(f"Error saving filtered JSON data: {e}")
    exit(1)

# Print results
print(f"Projection: Select only name and salary columns")
print(f"{projected_df}")
```

```
print(f"Aggregation: Calculate total salary")
```

```
print(f"Total Salary: {total_salary}")
```

```
print(f"\n")
```

```
print(f"# Count: Number of employees earning more than 50000")
```

```
print(f"Number of High Earners (>50000): {high_earners_count}")
```

```
print(f"\n")
```

```
print(f"limit Top 5 highest salary")
```

```
print(f"Top 5 Earners: \n{top_5_earners}")
```

```
print(f"\n")
```

```
print(f"Skipped DataFrame (First 2 rows skipped): \n{skipped_df}")
```

```
print(f"\n")
```

```
print(f"Filtered DataFrame (Sales department removed): \n{filtered_df}")
```

run the file by

```
bash: python3 process_data.py
```

output

Filtered JSON file saved successfully.

Projection: Select only name and salary columns

	name	salary
0	John Doe	50000
1	Jane Smith	60000
2	Alice Johnson	70000
3	Bob Brown	55000
4	Charlie Black	80000

Aggregation: Calculate total salary

Total Salary: 315000

# Count: Number of employees earning more than 50000

Number of High Earners (>50000): 4

limit Top 5 highest salary

Top 5 Earners:

	name	age	department	salary
4	Charlie Black	45	IT	80000
2	Alice Johnson	35	Finance	70000
1	Jane Smith	25	IT	60000
3	Bob Brown	28	Marketing	55000
0	John Doe	30	HR	50000

Skipped DataFrame (First 2 rows skipped):

	name	age	department	salary
2	Alice Johnson	35	Finance	70000
3	Bob Brown	28	Marketing	55000
4	Charlie Black	45	IT	80000

Filtered DataFrame (Sales department removed):

	name	age	department	salary
0	John Doe	30	HR	50000
2	Alice Johnson	35	Finance	70000
3	Bob Brown	28	Marketing	55000